

# Data Science Assignment 1: Ritika Thakur (2022408) | Swarnima Prasad (2022525)

```
In [32]: import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import numpy as np
```

## Loading the data

```
In [33]: df = pd.read_csv('AutoMPG.csv')
```

```
In [34]: print(df.head())

print("\nData Information:")
print(df.info())

print("\nSummary Statistics:")
print(df.describe())

print("\nMissing Values:")
print(df.isnull().sum())
```

	mpg	cylinders	displacement	horsepower	weight	acceleration	model year \
0	18.0	8	307.0	130	3504.0	12.0	70.0
1	15.0	8	350.0	165	3693.0	11.5	70.0
2	18.0	8	318.0	150	3436.0	11.0	70.0
3	16.0	8	304.0	150	3433.0	12.0	70.0
4	17.0	8	302.0	140	3449.0	10.5	70.0

	origin
0	1.0
1	1.0
2	1.0
3	1.0
4	1.0

Data Information:

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 399 entries, 0 to 398

Data columns (total 8 columns):

#	Column	Non-Null Count	Dtype
0	mpg	398 non-null	float64
1	cylinders	399 non-null	int64
2	displacement	398 non-null	float64
3	horsepower	398 non-null	object
4	weight	398 non-null	float64
5	acceleration	398 non-null	float64
6	model year	398 non-null	float64
7	origin	398 non-null	float64

dtypes: float64(6), int64(1), object(1)

memory usage: 25.1+ KB

None

Summary Statistics:

	mpg	cylinders	displacement	weight	acceleration \
count	398.000000	399.000000	398.000000	398.000000	398.000000
mean	23.514573	5.461153	193.425879	2960.047739	15.568090
std	7.815984	1.703638	104.269838	882.516758	2.757689
min	9.000000	3.000000	68.000000	-2065.000000	8.000000
25%	17.500000	4.000000	104.250000	2223.750000	13.825000
50%	23.000000	4.000000	148.500000	2803.500000	15.500000
75%	29.000000	8.000000	262.000000	3608.000000	17.175000
max	46.600000	8.000000	455.000000	5140.000000	24.800000

	model year	origin
count	398.000000	398.000000
mean	327.062814	1.572864
std	5008.688771	0.802055
min	70.000000	1.000000
25%	73.000000	1.000000
50%	76.000000	1.000000
75%	79.000000	2.000000
max	99999.000000	3.000000

Missing Values:

mpg	1
cylinders	0
displacement	1
horsepower	1
weight	1
acceleration	1

```
model year    1
origin        1
dtype: int64
```

Dropping the redundant last row

```
In [35]: # Drop the last row using the index of the last row
df.drop(df.index[-1], inplace=True)
```

## Handling Missing Values

```
In [36]: # Replace '?' with NaN
df['horsepower'].replace('?', pd.NA, inplace=True)

# Drop rows where 'horsepower' is missing
df.dropna(subset=['horsepower'], inplace=True)

# Convert 'horsepower' to numeric after dropping rows
df['horsepower'] = pd.to_numeric(df['horsepower'])

# Verify the changes
print(df['horsepower'].head())
```

```
0    130
1    165
2    150
3    150
4    140
```

Name: horsepower, dtype: int64

C:\Users\Ritika\AppData\Local\Temp\ipykernel\_52084\2527508700.py:2: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.

The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or 'df[col] = df[col].method(value)' instead, to perform the operation inplace on the original object.

```
df['horsepower'].replace('?', pd.NA, inplace=True)
```

```
In [37]: # Verify the changes
print(df.tail())
print("\nMissing Values:")
print(df.isnull().sum())
```

	mpg	cylinders	displacement	horsepower	weight	acceleration	\
393	27.0	4	140.0	86	2790.0	15.6	
394	44.0	4	97.0	52	2130.0	24.6	
395	32.0	4	135.0	84	2295.0	11.6	
396	28.0	4	120.0	79	2625.0	18.6	
397	31.0	4	119.0	82	2720.0	19.4	

	model year	origin
393	82.0	1.0
394	82.0	2.0
395	82.0	1.0
396	82.0	1.0
397	82.0	1.0

Missing Values:

```
mpg          0
cylinders    0
displacement 0
horsepower   0
weight       0
acceleration 0
model year   0
origin       0
dtype: int64
```

In [38]: `print(df.info())`

```
<class 'pandas.core.frame.DataFrame'>
Index: 392 entries, 0 to 397
Data columns (total 8 columns):
#   Column          Non-Null Count  Dtype
---  ---
0   mpg             392 non-null   float64
1   cylinders        392 non-null   int64
2   displacement     392 non-null   float64
3   horsepower       392 non-null   int64
4   weight           392 non-null   float64
5   acceleration     392 non-null   float64
6   model year      392 non-null   float64
7   origin           392 non-null   float64
dtypes: float64(6), int64(2)
memory usage: 27.6 KB
None
```

```
In [39]: # Total number of rows before handling '?'
initial_row_count = pd.read_csv('AutoMPG.csv').shape[0]

# Total number of rows after handling '?'
final_row_count = df.shape[0]

print(f"Initial number of rows: {initial_row_count}")
print(f"Final number of rows: {final_row_count}")
```

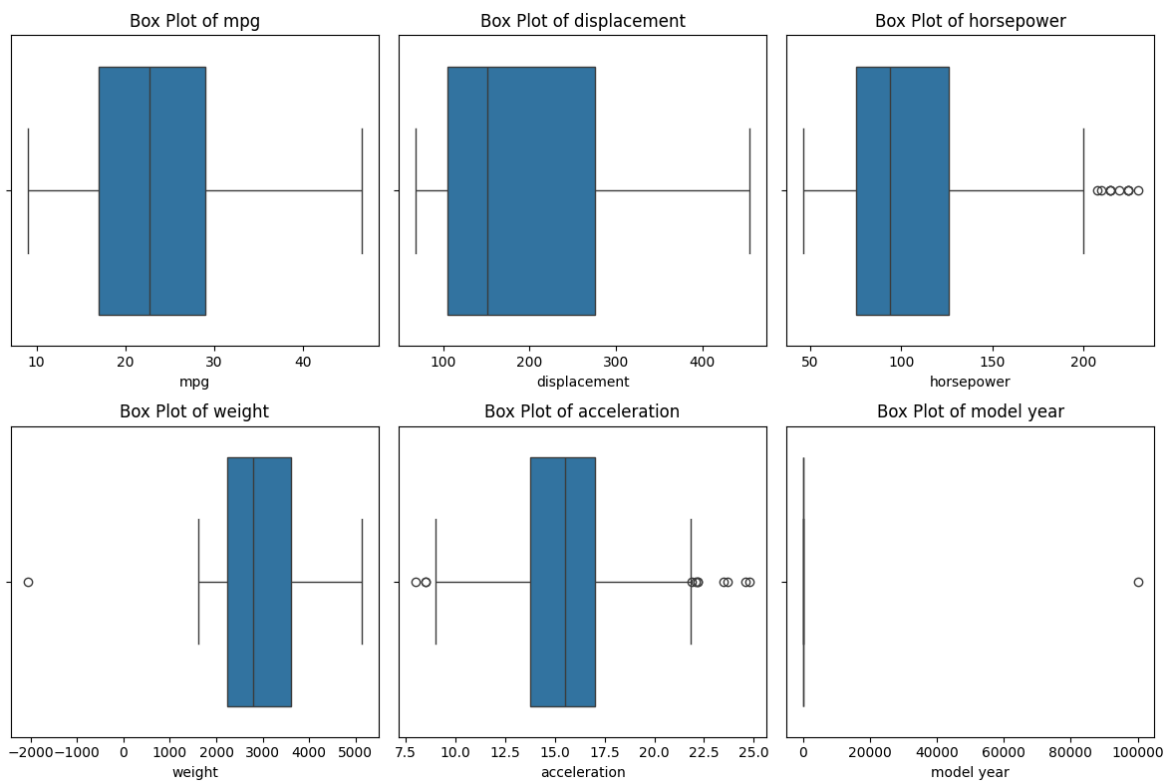
Initial number of rows: 399

Final number of rows: 392

## Trying outlier detection Techniques

In [40]: `numerical_columns = ['mpg', 'displacement', 'horsepower', 'weight', 'acceleratio`

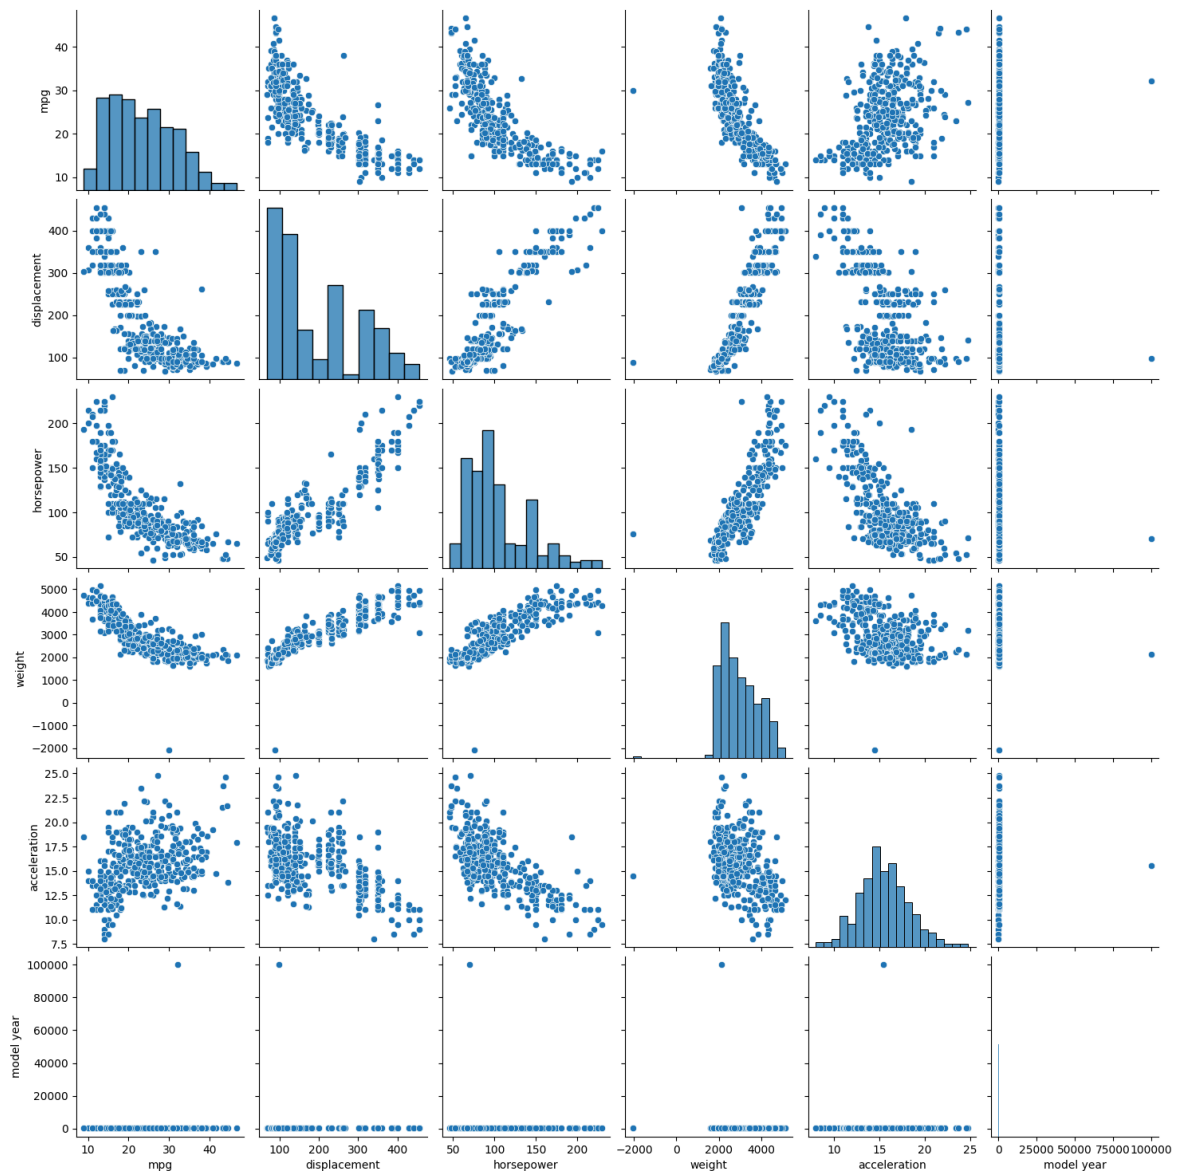
```
plt.figure(figsize=(12, 8))
for i, col in enumerate(numerical_columns, 1):
    plt.subplot(2, 3, i)
    sns.boxplot(x=df[col])
    plt.title(f'Box Plot of {col}')
plt.tight_layout()
plt.show()
```



Horsepower and acceleration are the two features with most outliers.

```
In [41]: # Scatter plots to check relationships and outliers
plt.figure(figsize=(12, 8))
sns.pairplot(df[numerical_columns])
plt.show()
```

<Figure size 1200x800 with 0 Axes>



```
In [42]: from scipy.stats import zscore

# Calculate Z-scores for numerical columns
z_scores = df[numerical_columns].apply(zscore)

outliers = (z_scores.abs() > 3).sum()
print("Number of outliers in each column based on Z-score:")
print(outliers)
```

Number of outliers in each column based on Z-score:

```
mpg          0
displacement 0
horsepower   5
weight       1
acceleration 2
model year   1
dtype: int64
```

Outliers with Z-score > 3 or < -3

```
In [43]: # Function to identify outliers using IQR
def detect_outliers_iqr(df, col):
    Q1 = df[col].quantile(0.25)
    Q3 = df[col].quantile(0.75)
```

```

IQR = Q3 - Q1
lower_bound = Q1 - 1.5 * IQR
upper_bound = Q3 + 1.5 * IQR
return df[(df[col] < lower_bound) | (df[col] > upper_bound)]

# Check outliers for each numerical column
for col in numerical_columns:
    outliers_iqr = detect_outliers_iqr(df, col)
    print(f'Number of outliers in {col} based on IQR: {len(outliers_iqr)}')

```

Number of outliers in mpg based on IQR: 0  
 Number of outliers in displacement based on IQR: 0  
 Number of outliers in horsepower based on IQR: 10  
 Number of outliers in weight based on IQR: 1  
 Number of outliers in acceleration based on IQR: 11  
 Number of outliers in model year based on IQR: 1

## Removing Outliers

```

In [44]: # Function to calculate outlier bounds
def calculate_outlier_bounds(column):
    Q1 = df[column].quantile(0.25)
    Q3 = df[column].quantile(0.75)
    IQR = Q3 - Q1
    lower_bound = Q1 - 1.5 * IQR
    upper_bound = Q3 + 1.5 * IQR
    return lower_bound, upper_bound

# Calculate outlier bounds for each numerical column
initial_shape = df.shape
outlier_bounds = {col: calculate_outlier_bounds(col) for col in numerical_columns}
print(outlier_bounds)

```

```

{'mpg': (np.float64(-1.0), np.float64(47.0)), 'displacement': (np.float64(-151.125), np.float64(531.875)), 'horsepower': (np.float64(-1.5), np.float64(202.5)), 'weight': (np.float64(141.0), np.float64(5699.0)), 'acceleration': (np.float64(8.900000000000002), np.float64(21.899999999999995)), 'model year': (np.float64(64.0), np.float64(88.0))}

```

```

In [45]: # Function to remove outliers based on bounds
def remove_outliers(df, column, lower_bound, upper_bound):
    return df[(df[column] >= lower_bound) & (df[column] <= upper_bound)]

# Apply outlier removal for each numerical column
for col in numerical_columns:
    lower_bound, upper_bound = outlier_bounds[col]
    df = remove_outliers(df, col, lower_bound, upper_bound)

```

```

In [46]: print(f"Shape before removing outliers: {initial_shape}")
print(f"Shape after removing outliers: {df.shape}")

```

Shape before removing outliers: (392, 8)  
 Shape after removing outliers: (370, 8)

```

In [47]: # Save DataFrame to a CSV file
df.to_csv('Cleaned_AutoMPG.csv', index=False)

```

```

In [48]: df = pd.read_csv('Cleaned_AutoMPG.csv')

```

```
In [49]: print("\nCleaned Data Summary:")
print(df.describe())
```

Cleaned Data Summary:

	mpg	cylinders	displacement	horsepower	weight \
count	370.000000	370.000000	370.000000	370.000000	370.000000
mean	23.612703	5.418919	189.458108	102.035135	2948.648649
std	7.633255	1.674022	98.730198	33.684675	829.145929
min	9.000000	3.000000	68.000000	46.000000	1613.000000
25%	17.600000	4.000000	105.000000	76.000000	2223.750000
50%	23.000000	4.000000	146.000000	92.500000	2781.500000
75%	29.000000	6.000000	258.000000	120.000000	3556.000000
max	46.600000	8.000000	429.000000	200.000000	5140.000000

	acceleration	model year	origin
count	370.000000	370.000000	370.000000
mean	15.549189	76.113514	1.594595
std	2.408187	3.622884	0.818468
min	9.500000	70.000000	1.000000
25%	14.000000	73.000000	1.000000
50%	15.500000	76.000000	1.000000
75%	17.000000	79.000000	2.000000
max	21.800000	82.000000	3.000000

	mpg	cylinders	displacement	horsepower	weight \
count	370.000000	370.000000	370.000000	370.000000	370.000000
mean	23.612703	5.418919	189.458108	102.035135	2948.648649
std	7.633255	1.674022	98.730198	33.684675	829.145929
min	9.000000	3.000000	68.000000	46.000000	1613.000000
25%	17.600000	4.000000	105.000000	76.000000	2223.750000
50%	23.000000	4.000000	146.000000	92.500000	2781.500000
75%	29.000000	6.000000	258.000000	120.000000	3556.000000
max	46.600000	8.000000	429.000000	200.000000	5140.000000

	acceleration	model year	origin
count	370.000000	370.000000	370.000000
mean	15.549189	76.113514	1.594595
std	2.408187	3.622884	0.818468
min	9.500000	70.000000	1.000000
25%	14.000000	73.000000	1.000000
50%	15.500000	76.000000	1.000000
75%	17.000000	79.000000	2.000000
max	21.800000	82.000000	3.000000


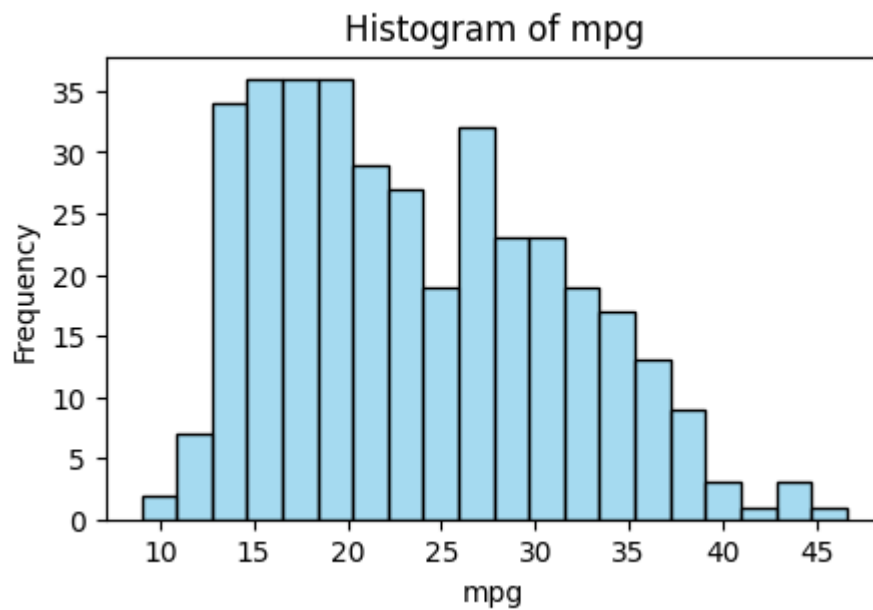
## Plotting histograms to assess the distribution of the data




```
In [50]: def plot_histogram(attribute):
            if attribute in df.columns:
                plt.figure(figsize=(5, 3))
                sns.histplot(df[attribute], kde=False, bins=20, color="skyblue")
                plt.title(f'Histogram of {attribute}')
                plt.xlabel(attribute)
                plt.ylabel('Frequency')
                plt.show()
            else:
                print(f"{attribute} not found in dataset.")

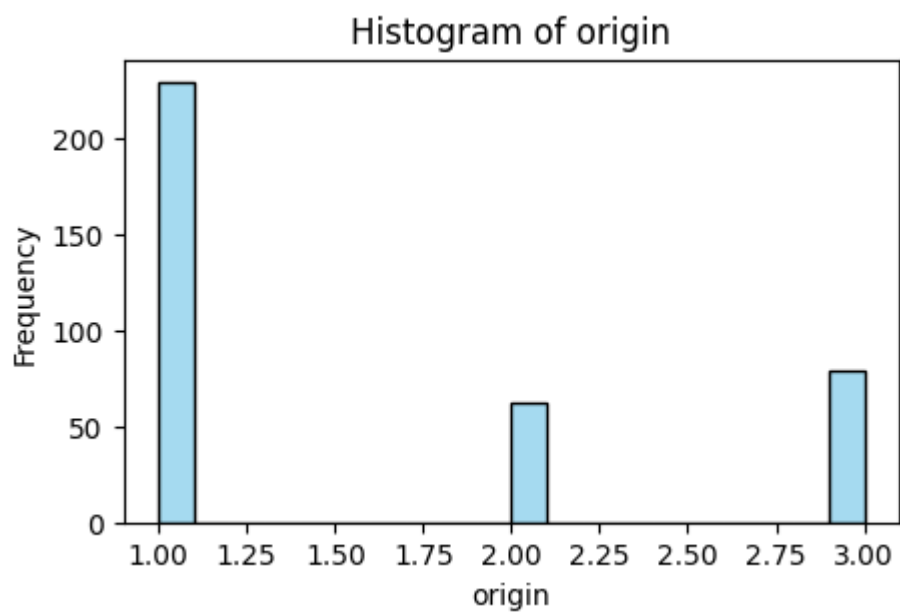


            # Plot histograms for continuous attributes
            plot_histogram('mpg')
            plot_histogram('displacement')
```



```
plot_histogram('weight')  
plot_histogram('acceleration')  
plot_histogram('model year')  
plot_histogram('horsepower')  
plot_histogram('cylinders')  
plot_histogram('origin')
```







```
In [51]: features = df.select_dtypes(include=[np.number])
print(features)
```

	mpg	cylinders	displacement	horsepower	weight	acceleration	\
0	18.0	8	307.0	130	3504.0	12.0	
1	15.0	8	350.0	165	3693.0	11.5	
2	18.0	8	318.0	150	3436.0	11.0	
3	16.0	8	304.0	150	3433.0	12.0	
4	17.0	8	302.0	140	3449.0	10.5	
..	...	...	...	...	...	...	
365	27.0	4	151.0	90	2950.0	17.3	
366	27.0	4	140.0	86	2790.0	15.6	
367	32.0	4	135.0	84	2295.0	11.6	
368	28.0	4	120.0	79	2625.0	18.6	
369	31.0	4	119.0	82	2720.0	19.4	

	model	year	origin
0		70.0	1.0
1		70.0	1.0
2		70.0	1.0
3		70.0	1.0
4		70.0	1.0
..		...	...
365		82.0	1.0
366		82.0	1.0
367		82.0	1.0
368		82.0	1.0
369		82.0	1.0

[370 rows x 8 columns]

## Calculating Mean , Variance and Squared Deviation on attributes

### Mean of each feature

$$\mu = \frac{\sum_{i=1}^n x_i}{n}$$

```
In [52]: def calculate_mean(df):
num_rows = len(df)
mean_vector = df.sum() / num_rows
return mean_vector

mean_vector = calculate_mean(df)
print("Mean of each feature:\n", mean_vector)
```

```
Mean of each feature:
mpg          23.612703
cylinders     5.418919
displacement 189.458108
horsepower   102.035135
weight      2948.648649
acceleration  15.549189
model year   76.113514
origin        1.594595
dtype: float64
```

### Variance of each feature

$$\sigma^2 = \frac{\sum_{i=1}^n (x_i - \mu)^2}{n}$$

```
In [53]: def calculate_variance(df, mean_vector):
    num_rows = len(df)
    variance = ((df - mean_vector) ** 2).sum() / num_rows # Manual calculation
    return variance

variance_vector = calculate_variance(df, mean_vector)
print("Variance of each feature:\n", variance_vector)
```

Variance of each feature:

```
mpg          58.109109
cylinders     2.794777
displacement  9721.307029
horsepower    1131.590657
weight       685624.908985
acceleration   5.783689
model year    13.089817
origin        0.668079
dtype: float64
```

## Total Variance

$$\sigma^2 = \frac{1}{n} \sum_{i=1}^n (x_i - \mu)^T (x_i - \mu)$$

```
In [54]: total_variance = 0
for i in range(len(df)):
    total_variance += np.dot((df.iloc[i] - mean_vector).T, (df.iloc[i] - mean_vector))
total_variance = total_variance / len(df)
print("\nTotal Variance ( $\sigma^2$ ):")
print(total_variance)
```

```
Total Variance ( $\sigma^2$ ):
696558.2521418552
```

## Normalizing Data

Standard deviation:

$$\sigma = \sqrt{\frac{\sum_{i=1}^n (x_i - \mu)^2}{n}}$$

Standardization:

$$x_{new} = \frac{x - \mu}{\sigma}$$

```
In [55]: def standardize_data(df):
    mean_vector = df.mean() # Compute the mean for each column
    print(mean_vector)
    std_vector = (calculate_variance(df, mean_vector))**0.5 # Compute the standard deviation
    print(std_vector)
    standardized_df = (df - mean_vector) / std_vector # Perform standardization
    return standardized_df, mean_vector, std_vector
```

```
# Normalize the cleaned dataset
standardized_df, mean_vector, std_vector = standardize_data(df)
print("Standardized data:\n", standardized_df)
```

```
mpg          23.612703
cylinders     5.418919
displacement 189.458108
horsepower   102.035135
weight       2948.648649
acceleration  15.549189
model year    76.113514
origin        1.594595
dtype: float64
mpg          7.622933
cylinders     1.671759
displacement  98.596689
horsepower    33.639124
weight       828.024703
acceleration   2.404930
model year     3.617985
origin         0.817361
dtype: float64
Standardized data:
      mpg  cylinders  displacement  horsepower   weight  acceleration \
0  -0.736292   1.543932    1.192148    0.831320  0.670694   -1.475797
1  -1.129841   1.543932    1.628269    1.871775  0.898948   -1.683704
2  -0.736292   1.543932    1.303714    1.425865  0.588571   -1.891610
3  -0.998658   1.543932    1.161721    1.425865  0.584948   -1.475797
4  -0.867475   1.543932    1.141437    1.128593  0.604271   -2.099516
..      ...      ...      ...      ...      ...      ...
365  0.444356  -0.848758   -0.390055   -0.357772  0.001632    0.728009
366  0.444356  -0.848758   -0.501620   -0.476681 -0.191599    0.021128
367  1.100272  -0.848758   -0.552332   -0.536136 -0.789407   -1.642122
368  0.575539  -0.848758   -0.704467   -0.684772 -0.390868    1.268565
369  0.969089  -0.848758   -0.714609   -0.595590 -0.276137    1.601215

      model year   origin
0    -1.689756 -0.727457
1    -1.689756 -0.727457
2    -1.689756 -0.727457
3    -1.689756 -0.727457
4    -1.689756 -0.727457
..      ...      ...
365    1.627007 -0.727457
366    1.627007 -0.727457
367    1.627007 -0.727457
368    1.627007 -0.727457
369    1.627007 -0.727457
```

[370 rows x 8 columns]

```
In [56]: sns.pairplot(standardized_df)
plt.show()
```



```
In [57]: mean_scaled = standardized_df.mean()
print(mean_scaled)
variance_scaled=calculate_variance(standardized_df, mean_scaled)
scaled_std_vector = (variance_scaled)**0.5
print(variance_scaled)
```

```
mpg          -3.840772e-16
cylinders    -1.728347e-16
displacement -1.344270e-16
horsepower   -5.761157e-17
weight       -1.536309e-16
acceleration  9.601929e-17
model year   9.217852e-16
origin       1.152231e-16
dtype: float64
mpg          1.0
cylinders    1.0
displacement 1.0
horsepower   1.0
weight       1.0
acceleration 1.0
model year   1.0
origin       1.0
dtype: float64
```

## Total Variance After normalizing

```
In [58]: total_variance = 0
for i in range(len(standardized_df)):
    total_variance += np.dot((standardized_df.iloc[i] - mean_scaled).T, (standardized_df.iloc[i] - mean_scaled))
total_variance = total_variance / len(standardized_df)
print("\nTotal Variance ( $\sigma^2$ ):")
print(total_variance)
```

Total Variance ( $\sigma^2$ ):  
8.0000000000000004

8.0000000000000004

Convert 'model year' and 'cylinders' to categorical data types

```
In [59]: df['model year'] = df['model year'].astype('category')
df['cylinders'] = df['cylinders'].astype('category')

print(df.dtypes)
```



```
mpg                float64
cylinders          category
displacement       float64
horsepower         int64
weight            float64
acceleration       float64
model year         category
origin            float64
dtype: object
```


```
In [60]: sns.countplot(data=df, x='cylinders')
plt.title("Count of Cylinders")
plt.xlabel("Cylinders")
plt.ylabel("Frequency")
plt.show()

sns.countplot(data=df, x='model year')
plt.title("Count of Model Years")
plt.xlabel("Model Year")
plt.ylabel("Frequency")
plt.show()

sns.countplot(data=df, x='model year', hue='cylinders')
plt.title("Cylinders by Model Year")
plt.xlabel("Model Year")
plt.ylabel("Count")
plt.show()
```







This visualization shows how the cylinders and model year values are categorical. Since both variables are categorical, the Chi-Square Test is appropriate for testing the independence of these variables.

## Chi-Square Test

Null Hypothesis ( $H_0$ ): There is no significant association between the number of cylinders and the model year of the cars in the dataset. In other words, the variables model year and cylinders are independent.

Alternative Hypothesis ( $H_1$ ): There is a significant association between the number of cylinders and the model year of the cars. In other words, the variables model year and cylinders are not independent.

Chi-Square:

$$\chi^2 = \sum \frac{(O_i - E_i)^2}{E_i}$$

```
In [61]: contingency_table = pd.crosstab(df['model year'], df['cylinders'])
print("Contingency Table:")
print(contingency_table)

total = contingency_table.to_numpy().sum()

expected_table = np.outer(contingency_table.sum(axis=1), contingency_table.sum(a
expected_table = pd.DataFrame(expected_table, index=contingency_table.index, col
```



- <https://www.medcalc.org/manual/statistical-tables.php>

```
In [49]: import numpy as np
import matplotlib.pyplot as plt
```

## Creating population of 1,00,000 points uniformly distributed between 0.01 and 1000

```
In [50]: # creating data consisting of 100000 points uniformly distributed between 0.01
data = np.zeros(100000)
for i in range(100000):
    data[i] = 0.01 * (i + 1)

print(data)
```

```
[1.0000e-02 2.0000e-02 3.0000e-02 ... 9.9998e+02 9.9999e+02 1.0000e+03]
```

## Mean and true variance

Mean:

$$\mu = \frac{1}{n} \sum_{i=1}^n x_i$$

True Variance:

$$\sigma^2 = \frac{1}{n} \sum_{i=1}^n (x_i - \mu)^2$$

```
In [51]: # mean
data_sum = np.sum(data)
data_mean = data_sum / 100000

print("mean:", data_mean)

# true variance
data_variance = np.sum((data - data_mean)**2) / 100000

print("true variance:", data_variance)
```

```
mean: 500.005
```

```
true variance: 83333.3333250001
```

## Computing s1\_squared, s2\_squared and s3\_squared for a sample of 50 points with replacement

s1\_squared:

$$s1^2 = \frac{1}{n-1} \sum_{i=1}^n (x_i - \bar{x})^2$$

s2\_squared:

$$s2^2 = \frac{1}{n} \sum_{i=1}^n (x_i - \bar{x})^2$$

s3\_squared:

$$s3^2 = \frac{n}{n-1} s2^2$$

```
In [52]: def compute_sample_variance(sample):
    sample_mean = np.sum(sample) / 50

    s1_sqd = sum((xi - sample_mean) ** 2 for xi in sample) / (51)
    s2_sqd = sum((xi - sample_mean) ** 2 for xi in sample) / 50
    s3_sqd = sum((xi - sample_mean) ** 2 for xi in sample) / (49)

    return s1_sqd, s2_sqd, s3_sqd
```

Average\_s1\_squared:

$$\frac{1}{m} \sum_{i=1}^m s1^2$$

Average\_s2\_squared:

$$\frac{1}{m} \sum_{i=1}^m s2^2$$

Average\_s3\_squared:

$$\frac{1}{m} \sum_{i=1}^m s3^2$$

```
In [53]: s1 = []
s2 = []
s3 = []

avg_s1 = []
avg_s2 = []
avg_s3 = []

itr = 1000

for _ in range(itr):
    sample = np.random.choice(data, 50, replace=True)
    s1_sqd, s2_sqd, s3_sqd = compute_sample_variance(sample)
    s1.append(s1_sqd)
    s2.append(s2_sqd)
    s3.append(s3_sqd)

    avg_s1.append(np.mean(s1))
    avg_s2.append(np.mean(s2))
    avg_s3.append(np.mean(s3))

# print("s1:", s1)
# print("s2:", s2)
# print("s3:", s3)

# print("avg_s1:", avg_s1)
# print("avg_s2:", avg_s2)
# print("avg_s3:", avg_s3)

plt.figure(figsize=(10, 6))


plt.scatter(range(itr), avg_s1, label = r'Avg $s_1^2$', s=10, alpha=0.6)
plt.axhline(y=data_variance, color='r', linestyle='--', label = 'True Variance')

plt.scatter(range(itr), avg_s2, label = r'Avg $s_2^2$', s=10, alpha=0.6)
# plt.axhline(y=data_variance, color='r', linestyle='--', label = 'True Variance')
```

```
plt.scatter(range(itr), avg_s3, label = r'Avg $s_3^2$', s=10, alpha=0.6)
# plt.axhline(y=data_variance, color='r', linestyle='-', Label = 'True Variance')

plt.title('Convergence of Sample Variances to True Variance')
plt.xlabel('Number of Iterations')
plt.ylabel('Variance')
plt.legend()

plt.show()
```



## Inferences

We notice that  $\text{r'Avg } s_3^2\text{'}$  reaches the true variance more quickly and frequently compared to the rest. This is because:

- the formula for  $s_1^2$  uses  $n + 1$  in the denominator, which tends to underestimate the variance. This makes  $s_1^2$  a biased estimator that is slightly biased downwards.
- the formula for  $s_2^2$  uses  $n$  in the denominator, which also results in a biased estimator but less than  $s_1^2$ .
- the formula for  $s_3^2$  uses  $n - 1$  in the denominator, which is the unbiased sample variance estimator. This formula compensates for the fact that the sample mean (used in calculating the variance) is based on the same data, thus giving a better approximation of the population variance.

$s_3^2$  is called an unbiased estimator because it is corrected for small sample sizes by dividing by  $n - 1$ . In statistics, dividing by  $n - 1$  is known as Bessel's correction, which accounts for the fact that the sample mean is less variable than the true population mean. This correction results in a more accurate estimate of the population variance when sampling randomly. Therefore, it tends to converge to the true variance more quickly and frequently compared to the other two estimators.

As more samples are taken, the law of large numbers ensures that all three sample variances will eventually converge to the true variance. However, for small sample sizes,  $s_3^2$  is preferred due to its unbiased nature and better approximation of the population variance.

## References

- [https://en.wikipedia.org/wiki/Bessel%27s\\_correction](https://en.wikipedia.org/wiki/Bessel%27s_correction)
- [https://en.wikipedia.org/wiki/Variance#Sample\\_variance](https://en.wikipedia.org/wiki/Variance#Sample_variance)

In [ ]:



# DSC Assignment 1 Question-3

September 20, 2024

## 1 Part (a)

(a) Let the die be unbiased with  $k$  faces. We want to find the expected number of rolls until the number  $\lfloor \sqrt{k} \rfloor$  appears on the upward face.

Since the die is unbiased, the probability of rolling any particular number is  $\frac{1}{k}$ .

Let  $X$  be the random variable representing the number of rolls needed to obtain  $\lfloor \sqrt{k} \rfloor$  on the upward face. The number of rolls follows a geometric distribution with success probability  $p = \frac{1}{k}$ .

The expected value for the geometric distribution is:

$$E(X) = \frac{1}{p} = \frac{1}{\frac{1}{k}} = k$$

Thus, the expected number of rolls to see the number  $\lfloor \sqrt{k} \rfloor$  is  $k$ .

## Part (B): Coupon Collector Problem

### General Formulation

In the general case of the **Coupon Collector Problem**, we have  $K$  distinct items (coupons, grades, etc.). The goal is to collect all  $K$  distinct items, and we are interested in the expected number of trials (or rolls, or papers) required to collect all of them at least once.

Let:

- $Y_i$  denote the number of trials until the  $i$ th new distinct item is obtained.
- $X_i = Y_{i+1} - Y_i$  represent the number of trials between obtaining the  $i$ th new distinct item and the  $(i + 1)$ th new distinct item.

We are interested in calculating  $\mathbb{E}[Y_K]$ , the expected number of trials to collect all  $K$  items.

### Expectation of $X_i$

Each  $X_i$  is a geometric random variable, and the probability of success (obtaining a new distinct item) decreases as we collect more distinct items. Specifically:

$$\mathbb{P}(\text{new distinct item on the } i\text{th trial}) = \frac{K - i}{K}$$

Therefore, the expected number of trials to get the next distinct item is:

$$\mathbb{E}[X_i] = \frac{K}{K - i}$$

for  $i = 0, 1, 2, \dots, K - 1$ .

### Total Expected Number of Trials

To find the total expected number of trials to collect all  $K$  distinct items, we sum the expectations of  $X_i$  over all possible  $i$ . Thus, we have:

$$\mathbb{E}[Y_K] = \sum_{i=0}^{K-1} \mathbb{E}[X_i] = \sum_{i=0}^{K-1} \frac{K}{K - i}$$

This simplifies to:

$$\mathbb{E}[Y_K] = K \cdot \sum_{i=1}^K \frac{1}{i}$$

## Harmonic Numbers

The sum  $\sum_{i=1}^K \frac{1}{i}$  is known as the  $K$ th *harmonic number*, denoted by  $H_K$ . Therefore, the expected number of trials to collect all  $K$  distinct items can also be written as:

$$\mathbb{E}[Y_K] = K \cdot H_K$$

For large values of  $K$ , the harmonic number  $H_K$  can be approximated as:

$$H_K \approx \ln(K) + \gamma$$

where  $\gamma \approx 0.5772$  is the Euler-Mascheroni constant.

Thus, for large  $K$ , the expected number of trials is approximately:

$$\mathbb{E}[Y_K] \approx K \cdot (\ln(K) + \gamma)$$

## Conclusion

In the general case of the coupon collector problem, the expected number of trials to collect all  $K$  distinct items grows logarithmically with  $K$ , following the harmonic number approximation. This result is significant in fields such as probability theory and combinatorics, where such "collection problems" frequently arise.

## Part (c)

### 1 Unequal probability while rolling dice

Let's break down the steps to calculate the expected number of rolls using inclusion-exclusion.

#### 1.1 Problem Statement

We have a  $k$ -faced die, where each face  $i$  has a probability  $p_i$  of landing on that face. The goal is to compute the expected number of rolls until **all faces** have been rolled at least once.

#### 1.2 Inclusion-Exclusion Formula

The inclusion-exclusion formula for this type of problem is:

$$E = \sum_i \frac{1}{p_i} - \sum_{i,j} \frac{1}{p_i + p_j} + \sum_{i,j,k} \frac{1}{p_i + p_j + p_k} - \cdots + (-1)^{n-1} \frac{1}{p_1 + p_2 + \cdots + p_n}$$

- **First term:** The expected number of rolls to see any single face (e.g.,  $P(i)$ ) is the inverse of the probability of seeing that face:  $\frac{1}{p_i}$ .
- **Second term:** The expected number of rolls to see two faces (e.g., face  $i$  and face  $j$ ) is  $\frac{1}{p_i + p_j}$ , but we need to subtract this because we've double-counted rolls that show both faces.
- **Third term:** We add back the cases where three faces have been counted together (e.g.,  $\frac{1}{p_i + p_j + p_k}$ ), and so on.
- This alternating sum accounts for all possible overlaps between the faces.

#### 1.3 Probabilities for the Geometric Die

A geometric die is a die where the probabilities decrease geometrically. For example, for a 3-sided geometric die:

$$P(1) = \frac{1}{4}, \quad P(2) = \frac{1}{2}, \quad P(3) = \frac{1}{4}$$

More generally, for a  $k$ -sided geometric die, the probability of face  $i$  is  $P(i) = \frac{1}{2^{i-1}}$ . So the probability for the first face is  $\frac{1}{1}$ , the second face is  $\frac{1}{2}$ , the third is  $\frac{1}{4}$ , and so on.

## 1.4 Apply the Formula

For a 3-sided die with probabilities  $\frac{1}{4}, \frac{1}{2}, \frac{1}{4}$ , the expected number of rolls  $E$  is computed as:

$$E = \left( \frac{1}{P(1)} + \frac{1}{P(2)} + \frac{1}{P(3)} \right) - \left( \frac{1}{P(1) + P(2)} + \frac{1}{P(1) + P(3)} + \frac{1}{P(2) + P(3)} \right) + \frac{1}{P(1) + P(2) + P(3)}$$

Generalized formula:

$$\mathbb{E} = \sum_{\mu \neq 0} \frac{(-1)^{|\mu|-1}}{\mu \cdot \mathbf{p}},$$

Substituting the values:

$$E = (4 + 2 + 4) - \left( \frac{4}{3} + 2 + \frac{4}{3} \right) + 1$$

Simplifying:

$$E = 10 - 4.67 + 1 = 6.33$$

So the expected number of rolls to see all three faces at least once is approximately 6.33 rolls.

```
In [3]: import numpy as np
import matplotlib.pyplot as plt
from itertools import combinations
```

### Inclusion-Exclusion Principle

$$E[\text{rolls}] = \sum_{r=1}^n (-1)^{r-1} \sum_{S \subseteq \text{probs}, |S|=r} \frac{1}{\sum_{i \in S} p_i}$$

```
In [4]: def expected_rolls(probs):
n = len(probs)
expected_value = 0

for r in range(1, n+1):
    for subset in combinations(probs, r):
        prob_sum = sum(subset)
        expected_value += (-1)**(r - 1) * 1 / prob_sum

return expected_value
```

### Calculating probabilities for a k-faced geometric die

$$P(i) = \frac{1}{2^{(i-1)}} \quad \text{for } i = 1, 2, \dots, k$$

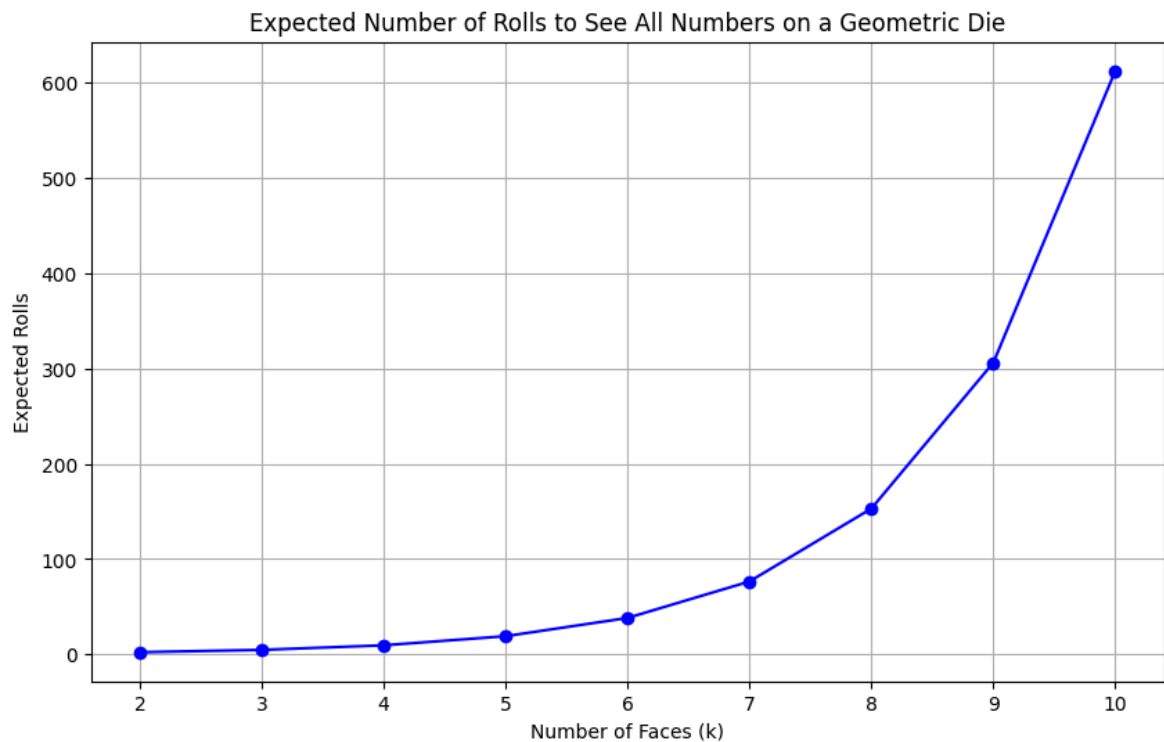
```
In [5]: def geometric_probabilities(k):
return [1/(2**(i-1)) for i in range(1, k+1)]
```

### Plotting the expected number of rolls for different values of k

```
In [6]: ks = range(2, 11)
expected_values = []

for k in ks:
    probabilities = geometric_probabilities(k)
    e_value = expected_rolls(probabilities)
    expected_values.append(e_value)

plt.figure(figsize=(10, 6))
plt.plot(ks, expected_values, marker='o', linestyle='--', color='b')
plt.title('Expected Number of Rolls to See All Numbers on a Geometric Die')
plt.xlabel('Number of Faces (k)')
plt.ylabel('Expected Rolls')
plt.grid(True)
plt.show()
```



The code simulates the expected number of rolls needed to observe all faces of a geometric die, where each face has a different probability of showing up. Specifically, the die has ( k ) faces, and the probability of each face being rolled follows a geometric distribution where the probability of rolling face ( i ) is

$$\frac{1}{2^{i-1}}$$

## Inference from the Plot:

### 1. Increasing ( k ) Increases Expected Rolls:

- As ( k ) (the number of faces of the die) increases, the expected number of rolls to see all faces at least once also increases because with more faces, it becomes harder to roll each unique number due to the geometric nature of the probabilities.

### 2. Geometric Distribution Effect:

- The probabilities of rolling each face decrease exponentially as ( i ) increases (the first face has a probability of ( 1 ), the second face has a probability of (  $\frac{1}{2}$  ), the third face has a probability of (  $\frac{1}{4}$  )...). This means that the later faces become progressively harder to roll, contributing significantly to the expected number of rolls required to observe all faces.

### 3. Diminishing Return of Rolls:

- The increase in expected rolls shows a slowing growth as ( k ) increases. This diminishing rate of increase can be attributed to the geometric probabilities: while adding more faces makes it harder to observe all of them, the difference between adding, say, the 9th and 10th faces (which have very low probabilities) is less significant than adding earlier faces (which have higher probabilities).

#### 4. Inclusion-Exclusion Principle:

- The code uses the inclusion-exclusion principle to calculate the expected number of rolls. This accounts for overlapping probabilities of combinations of faces, providing a more accurate estimate than a simple sum of probabilities would.

**Overall, difficulty of observing all numbers on a geometric die increases as the number of faces increases, with a slowing growth due to the nature of the geometric distribution.**

## References

- <https://www.jstor.org/stable/40378689?seq=3>
- <https://www.jstor.org/stable/40378689?seq=2>
- <https://www.youtube.com/watch?v=3mu47FWEuqA>
- <https://math.stackexchange.com/questions/600012/coupon-collectors-problem-with-unequal-probabilities>



```
In [135... import pandas as pd
import math
from scipy import stats
import matplotlib.pyplot as plt
```

## Loading the dataset

```
In [135... data = pd.read_csv('Hurricane.csv')
```

```
In [135... print(data.head())
```

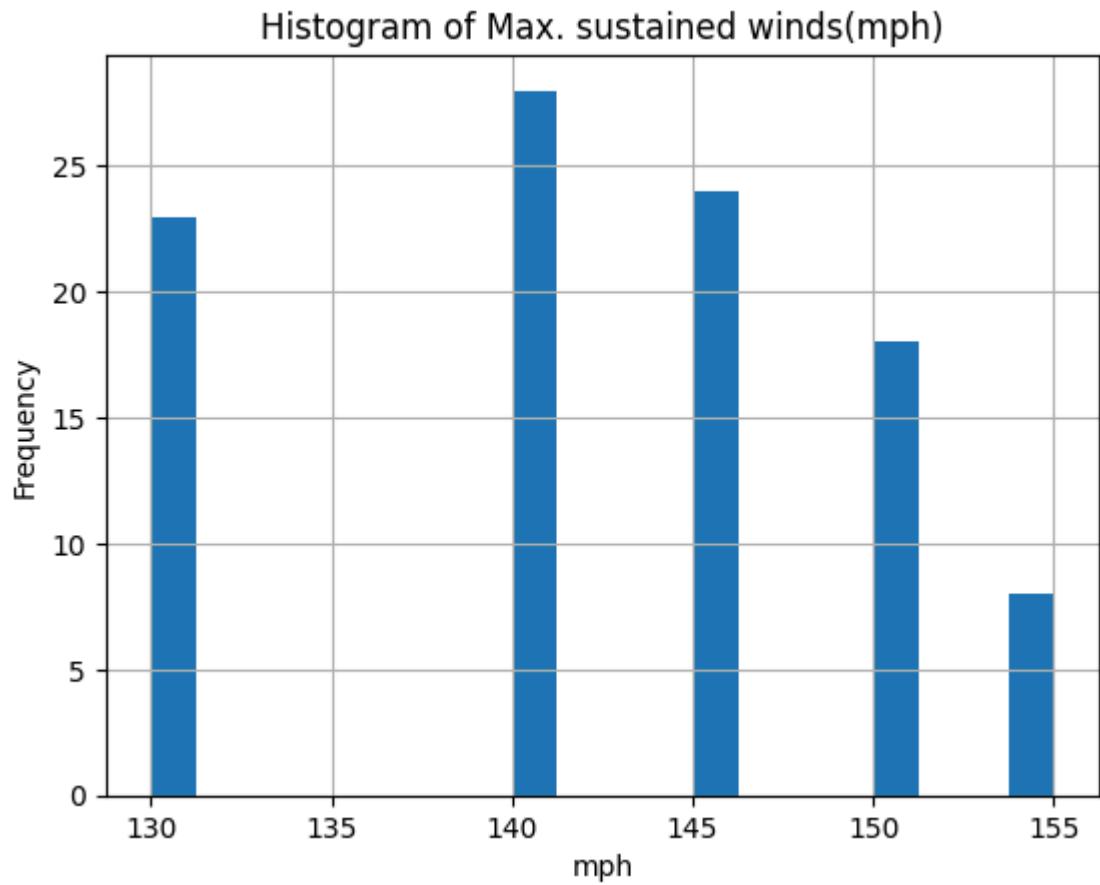
	Name	Season	Month \
0	Hurricane #3	1853	August, September
1	"1856 Last Island Hurricane"	1856	August
2	Hurricane #6	1866	September, October
3	Hurricane #7	1878	September, October
4	Hurricane #2	1880	August

	Max. sustained winds(mph)	Minimum pressure(mbar)
0	150	924
1	150	934
2	140	938
3	140	938
4	150	931


```
In [135... data.columns = ['Name', 'Season', 'Month', 'mph', 'mbar']
```

```
In [135... data['mph'].hist(bins=20)
plt.xlabel('mph')
plt.ylabel('Frequency')
plt.title('Histogram of Max. sustained winds(mph)')
plt.show()
```



In [136...

```
data['mbar'].hist(bins=20)
plt.xlabel('mbar')
plt.ylabel('Frequency')
plt.title('Histogram of Min. pressure(mbar)')
plt.show()
```



## Preprocessing

In [136...

```
# cleaning Month column
data['Month'] = data['Month'].str.strip()
data['Month'] = data['Month'].str.replace(',', ' ')
data['Month'] = data['Month'].str.replace('-', ' ')
data['Month'] = data['Month'].str.replace('Aug', 'August')
data['Month'] = data['Month'].str.replace(r'\d+', '', regex=True)
data['Month'] = data['Month'].str.replace('Augustust', 'August')

data['Month']
```

Out[136...

```
0      August  September
1              August
2      September  October
3      September  October
4              August
...
96              October
97      August  September
98              August
99      August  September
100             September
Name: Month, Length: 101, dtype: object
```

In [136...

```
# splitting one row into multiple rows if it contains multiple months
data['Month'] = data['Month'].str.split()
data = data.explode('Month')

data
```

Out[136...

	Name	Season	Month	mph	mbar
0	Hurricane #3	1853	August	150	924
0	Hurricane #3	1853	September	150	924
1	"1856 Last Island Hurricane"	1856	August	150	934
2	Hurricane #6	1866	September	140	938
2	Hurricane #6	1866	October	140	938
...	...	...	...	...	...
97	Hurricane Fabian	2003	September	145	939
98	Hurricane Charley	2004	August	150	941
99	Hurricane Frances	2004	August	145	935
99	Hurricane Frances	2004	September	145	935
100	Hurricane Karl	2004	September	145	938

137 rows × 5 columns

(a) With a 1% level of significance conduct t-test for correlation coefficient between "Max. sustained winds(mph)" and "Minimum pressure(mbar)".

### Preprocessing the data

In [136...

```
# min max scaling numerical data
# mph_org = data['mph']
# mbar_org = data['mbar']
# data['mph'] = (data['mph'] - data['mph'].min()) / (data['mph'].max() - data['m
# data['mbar'] = (data['mbar'] - data['mbar'].min()) / (data['mbar'].max() - dat
```

### t-test

Null Hypothesis: There is no correlation between "Max. sustained winds(mph)" and "Minimum pressure(mbar)"

Alternate Hypothesis: There is a correlation between "Max. sustained winds(mph)" and "Minimum pressure(mbar)"

Method:\

Covariance:\

$$Cov(X, Y) = \frac{\sum_{i=1}^n (X_i - \bar{X})(Y_i - \bar{Y})}{n}$$

Standard Deviation:\

$$\sigma_X = \sqrt{\frac{\sum_{i=1}^n (X_i - \bar{X})^2}{n}}$$

Correlation Coefficient:\

$$r = \frac{Cov(X, Y)}{\sigma_X \sigma_Y}$$

t-test:\

$$t = \frac{r\sqrt{n-2}}{\sqrt{1-r^2}}$$

where n is the number of samples

In [136...

```
# Calculating mean
mean_mph = data['mph'].mean()
mean_mbar = data['mbar'].mean()

print(f"Mean of Max Sustained Winds: {mean_mph}")
print(f"Mean of Minimum Pressure: {mean_mbar}")

# Calculating covariance
covariance = sum((data['mph'] - mean_mph) * (data['mbar'] - mean_mbar)) / (len(data) - 1)

print(f"Covariance: {covariance}")

# Calculating standard deviation
std_dev_mph = math.sqrt(sum((data['mph'] - mean_mph) ** 2) / (len(data) - 1))
std_dev_mbar = math.sqrt(sum((data['mbar'] - mean_mbar) ** 2) / (len(data) - 1))

print(f"Standard Deviation of Max Sustained Winds: {std_dev_mph}")
print(f"Standard Deviation of Minimum Pressure: {std_dev_mbar}")

# Calculating correlation
correlation = covariance / (std_dev_mph * std_dev_mbar)

print(f"Correlation: {correlation}")

# t-test
n = len(data)
t = correlation * math.sqrt(n - 2) / math.sqrt(1 - correlation ** 2)

print(f"t-test: {t}")

# range: mean - t * std_dev, mean + t * std_dev
t_value = stats.t.ppf(0.995, n-2)

print(f"t-value: {t_value}")

print("-----")

if abs(t) > t_value:
    print("Reject the null hypothesis; there is a significant correlation between")
else:
    print("Accept the null hypothesis; there is no significant correlation between")

print("-----")

# p-value
p = 2 * (1 - stats.t.cdf(abs(t), df=n-2))
```

```

print("Degree of Freedom:", n-2)

print(f"p-value: {p}")

print("-----")

# Conclusion
if p < 0.01:
    print("Reject the null hypothesis; there is a significant correlation between")
else:
    print("Accept the null hypothesis; there is no significant correlation between")

```

Mean of Max Sustained Winds: 142.33576642335765

Mean of Minimum Pressure: 938.8029197080292

Covariance: -35.89657578359814

Standard Deviation of Max Sustained Winds: 7.766138199470568

Standard Deviation of Minimum Pressure: 9.985165037413662

Correlation: -0.46290584088602793

t-test: -6.067728544364732

t-value: 2.612737907693308

-----

Reject the null hypothesis; there is a significant correlation between Max Sustained Winds and Minimum Pressure

-----

Degree of Freedom: 135

p-value: 1.2303904561861145e-08

-----

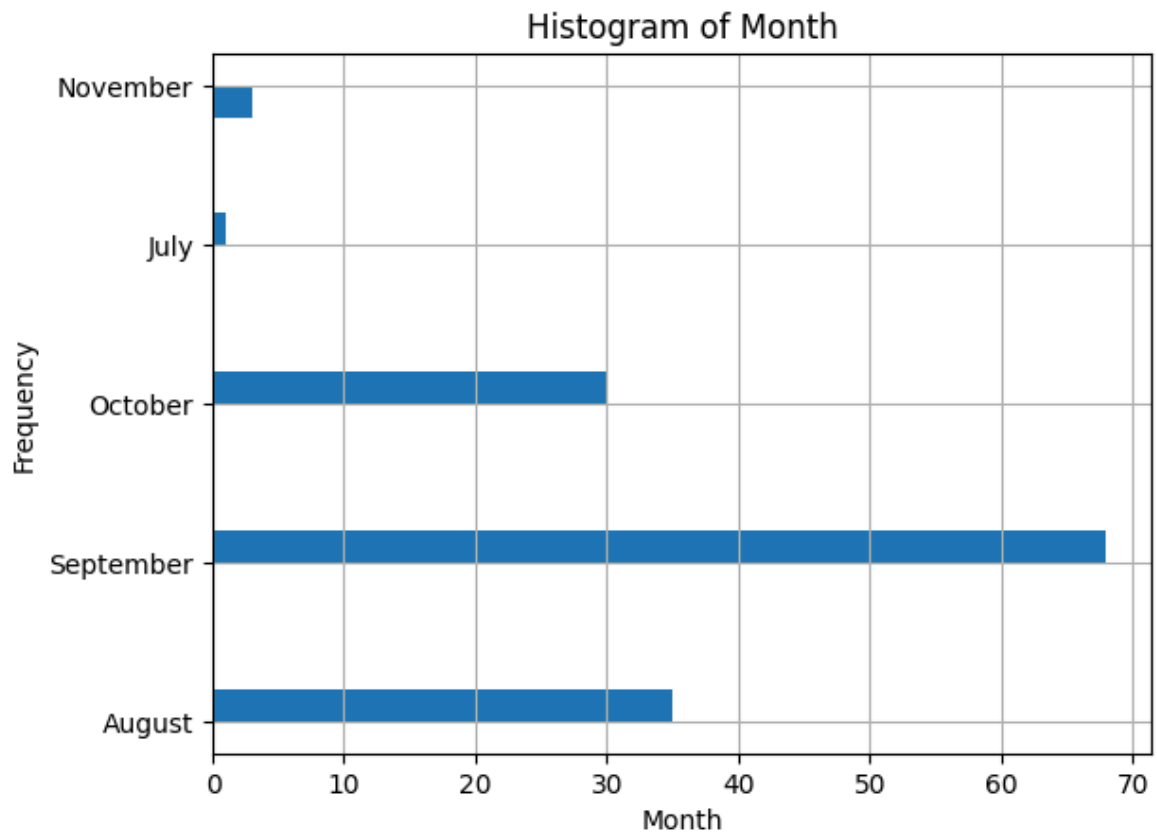
Reject the null hypothesis; there is a significant correlation between Max Sustained Winds and Minimum Pressure

**(b) With a 5% level of significance test if the "Max. sustained winds(mph)" of hurricane depends on the month of its occurrence.**

```


In [136... data['Month'].hist(bins=20, orientation='horizontal')
plt.xlabel('Month')
plt.ylabel('Frequency')
plt.title('Histogram of Month')
plt.show()

```



```
In [138...] contingency_table = pd.crosstab(data['Month'], data['mph'])
contingency_table.plot(kind='bar', stacked=True)
contingency_table
```

```
Out[138...]    mph  130  140  145  150  155
Month
0      0    1   0    0    0
1      7    9   7    8    4
2     16   18  17   11    6
3      5   10   8    5    2
4      0    1   1    0    1
```



### Chi Square test

Null Hypothesis: The "Max. sustained winds(mph)" of hurricane does not depend on the month of its occurrence

Alternate Hypothesis: The "Max. sustained winds(mph)" of hurricane depends on the month of its occurrence

Method:\

Chi Square test:\

$$\chi^2 = \sum_{i=1}^n \frac{(O_i - E_i)^2}{E_i}$$

where n is the number of categories

```
In [136... month_total = contingency_table.sum(axis=1)
mph_total = contingency_table.sum(axis=0)
contingency_total = contingency_table.values.sum()

print("month_total:", month_total)
print("mph_total:", mph_total)
print("contingency_total:", contingency_total)
```



```

month_total: Month
August      35
July        1
November    3
October     30
September   68
dtype: int64
mph_total: mph
130         28
140         39
145         33
150         24
155         13
dtype: int64
contingency_total: 137

```

```

In [136... expected_frequency = pd.DataFrame(index=contingency_table.index, columns=contingency_table.columns)

for i in contingency_table.index:
    for j in contingency_table.columns:
        expected_frequency.at[i, j] = (month_total[i] * mph_total[j]) / contingency_total

print("Expected Frequency:")
print(expected_frequency)

```

```

Expected Frequency:
mph      130      140      145      150      155
Month
August    7.153285  9.963504  8.430657  6.131387  3.321168
July       0.20438  0.284672  0.240876  0.175182  0.094891
November   0.613139  0.854015  0.722628  0.525547  0.284672
October    6.131387  8.540146  7.226277  5.255474  2.846715
September 13.89781 19.357664 16.379562 11.912409  6.452555

```

```

In [136... # Chi-square test
chi_square = 0

for i in contingency_table.index:
    for j in contingency_table.columns:
        observed = contingency_table.at[i, j]
        expected = expected_frequency.at[i, j]
        chi_square += (observed - expected) ** 2 / expected

print(f"Chi-square: {chi_square}")

```

Chi-square: 7.971634778620072

```

In [137... rows, cols = contingency_table.shape
df = (rows - 1) * (cols - 1)

print(f"Degrees of Freedom: {df}")

```

Degrees of Freedom: 16

```

In [137... chi_square_critical = stats.chi2.ppf(1 - 0.025, df)
chi_square_critical_lower = stats.chi2.ppf(0.025, df)

print(f"Chi-square Critical: {chi_square_critical}")
print(f"Chi-square Critical Lower: {chi_square_critical_lower}")

if chi_square > chi_square_critical or chi_square < chi_square_critical_lower:

```

```
print("Reject the null hypothesis; there is a significant relationship between
else:
print("Accept the null hypothesis; there is no significant relationship between
```

Chi-square Critical: 28.845350723404753

Chi-square Critical Lower: 6.907664353497004

Accept the null hypothesis; there is no significant relationship between Month and Max Sustained Winds

In [137...

```
# p-value
p = 1 - stats.chi2.cdf(chi_square, df=df)

print(f"p-value: {p}")

# Conclusion
if p < 0.05:
    print("Reject the null hypothesis; there is a significant association between
else:
    print("Accept the null hypothesis; there is no significant association between
```

p-value: 0.9497063344388978

Accept the null hypothesis; there is no significant association between Month and Max Sustained Winds

## Further processing of 'Month' column

In [137...

```
# converting months into relative ordinal values
data['Month'] = pd.Categorical(data['Month'], ordered=True, categories=['July',
data['Month'] = data['Month'].cat.codes

print(data['Month'])
```

```
0      1
0      2
1      1
2      2
2      3
..
97     2
98     1
99     1
99     2
100    2
```

Name: Month, Length: 137, dtype: int8

## t-test

Null Hypothesis: There is no correlation between "Max. sustained winds(mph)" and "Month"

Alternate Hypothesis: There is a correlation between "Max. sustained winds(mph)" and "Month"

In [137...

```
# t-test
# calculating mean
mean_month = data['Month'].mean()

print(f"Mean of Month: {mean_month}")

# calculating covariance
```

```

covariance = sum((data['Month'] - mean_month) * (data['mph'] - mean_mph)) / (len
print(f"Covariance: {covariance}")

# calculating standard deviation
std_dev_month = math.sqrt(sum((data['Month'] - mean_month) ** 2) / (len(data) -

print(f"Standard Deviation of Month: {std_dev_month}")

# calculating correlation
correlation = covariance / (std_dev_month * std_dev_mph)

print(f"Correlation: {correlation}")

# t-test
n = len(data)
t = correlation * math.sqrt(n - 2) / math.sqrt(1 - correlation ** 2)

print(f"t-test: {t}")

# range: mean - t * std_dev, mean + t * std_dev
t_value = stats.t.ppf(0.975, n-2)

print(f"t-value: {t_value}")

if abs(t) > t_value:
    print("Reject the null hypothesis; there is a significant correlation between
else:
    print("Accept the null hypothesis; there is no significant correlation between

# p-value
p = 2 * (1 - stats.t.cdf(abs(t), df=n-2))

print("Degree of Freedom:", n-2)

print(f"p-value: {p}")

# Conclusion
if p < 0.05:
    print("Reject the null hypothesis; there is a significant correlation between
else:
    print("Accept the null hypothesis; there is no significant correlation between

```

Mean of Month: 1.9927007299270072

Covariance: 0.05393945899527636

Standard Deviation of Month: 0.7717088597331507

Correlation: 0.009000113479535317

t-test: 0.1045761043854583

t-value: 1.977692277222804

Accept the null hypothesis; there is no significant correlation between Month and Max Sustained Winds

Degree of Freedom: 135

p-value: 0.9168673863071735

Accept the null hypothesis; there is no significant correlation between Month and Max Sustained Winds

Degree of Freedom: 135

p-value: 0.9168673863071735

Accept the null hypothesis; there is no significant correlation between Month and Max Sustained Winds

**With a 10% level of significance conduct test if “Max. sustained winds(mph)” follows a Poisson distribution.**

Poisson distribution is given by:

$$P(X = k) = \frac{e^{-\lambda} \lambda^k}{k!}$$

where k is the number of occurrences, and  $\lambda$  is the average number of occurrences

### Expected and Observed frequencies

```
In [137... # scaling_factor = 2

# # scaling mph
# data['mph'] = data['mph'] // scaling_factor

In [137... # scaling_factor = 2

# data['mph'] = data['mph'] * scaling_factor
# data['mph'] = data['mph'].astype(int)
# data['mph']

In [137... #min max scaling month
# data['Month'] = (data['Month'] - data['Month'].min()) / (data['Month'].max() -

In [137... mph_count = data['mph'].value_counts()

print(mph_count)

mph
140    39
145    33
130    28
150    24
155    13
Name: count, dtype: int64

In [137... # # how many times each wind occurred
# mph_count = mph_org.value_counts()

# print("mph_count:")
# print(mph_count)

In [138... mean_mph_scaled = data['mph'].mean()

In [138... # expected frequencies using Poisson distribution
expected_frequency = {}
n = len(data)

# print("Total Observations:", n)

for i in mph_count.index:
    poisson_prob = stats.poisson.pmf(i, mean_mph_scaled)
    # poisson_prob = (mean_mph_scaled ** i) * math.exp(-mean_mph_scaled) / math.
    expected_frequency[i] = poisson_prob * n + 1e-15
```

```
print("Expected Frequency:")
print(expected_frequency)
```

Expected Frequency:

{140: 4.52833796439542, 145: 4.425222015865378, 130: 2.7619462507085055, 150: 3.6416302840133272, 155: 2.5380588506680306}

## Chi Square test

Null Hypothesis: "Max. sustained winds(mph)" follows a Poisson distribution

Alternate Hypothesis: "Max. sustained winds(mph)" does not follow a Poisson distribution

```
In [138... chi_square = 0

for i in mph_count.index:
    observed = mph_count[i]
    expected = expected_frequency[i]
    chi_square += (observed - expected) ** 2 / expected

print(f"Chi-square: {chi_square}")
```

Chi-square: 834.484409173985

```
In [138... df = len(mph_count) - 2

print(f"Degrees of Freedom: {df}")
```

Degrees of Freedom: 3

```
In [138... chi_square_critical = stats.chi2.ppf(0.90, df)

print(f"Chi-square Critical: {chi_square_critical}")

if abs(chi_square) > chi_square_critical:
    print("Reject the null hypothesis; the distribution of Max Sustained Winds is not Poisson")
else:
    print("Accept the null hypothesis; the distribution of Max Sustained Winds is Poisson")
```

Chi-square Critical: 6.251388631170325

Reject the null hypothesis; the distribution of Max Sustained Winds is not Poisson

## References

- <https://www.medcalc.org/manual/statistical-tables.php>
- <https://www.statology.org/t-test-for-correlation/>
- [https://stats.libretexts.org/Bookshelves/Introductory\\_Statistics/Introductory\\_Statistics\\_1e](https://stats.libretexts.org/Bookshelves/Introductory_Statistics/Introductory_Statistics_1e)

## Note:

High Level discussions conducted with the group consisting of: Saksham Singh and Sidhartha Garg

In [ ]: